

Understanding Utilitarianism and its Variants

Fundamental question: which variant is the best representation of utilitarianism? What form best achieves the goals of utilitarianism as an ethical theory? Are other variants of utilitarianism meaningless because they collapse into Act utilitarianism?

Fundamental Principle of Utilitarianism

Characterising utilitarianism

- Criterion of rightness (explanatory goal): An informative account of which actions are correct and why
- Decision making procedure (Deliberative goal): An account useful for deciding what to do

Components of utilitarianism

- Maximising, consequentialist, egalitarian

Interpreting Mill's principle of utility

- An action is right if and only if it produces as much welfare, hedonistically conceived, overall – where everyone's welfare counts and counts equally – as any other action open to the agent
- Actions are right insofar as they increase happiness, and wrong insofar as they decrease it
 - Crisp suggests that Mill interprets our rightness and wrongness in terms of degrees and that these qualities can exist simultaneously in the same action
 - But Mill suggests that the true right action is that which **produces the greatest balance of pleasure**
- Consequentialist assumption
 - Actions are the immediate focus of evaluation; everything else is ancillary or a means to producing the optimal outcome
- Crisp reformulates Mill's utilitarianism in terms of histories
 - The best actions are those which feature in the best possible history
 - Note that Crisp says nothing about one-shot simultaneous 'games' or actions in that the individual utility maximising action often leads to the socially suboptimal outcome
 - Rather, his notion of histories is **sequential** in nature; we must choose actions that maximise utility in the long run
 - We may choose two individually utility maximising actions at different points in time which conflict with one another
 - Though the principle is the same: when choosing actions across a particular dimension (time/persons), we must choose the overall maximising set of actions
 - Note that this is distinct from Rule Utilitarianism: a set of actions that collectively maximises utility does not imply that these actions form rules. Instead, it is just a collection of individual actions made by

different people across time which maximises utility (i.e. Matthew lying at t1 could feature in the best possible history while a teacher telling the truth at t3 also features in the best possible history)

- So, the principle of utility sanctions a set of actions that collectively maximises overall utility, rather than individually utility-maximising actions
- Actualism and Probabilism debate (actual value or expected value)
 - The Rash Doctor: You have a serious medical condition, for which two treatments are available. One will leave you with a high welfare level of 50, the other with a fairly low level of 25. There is, however, only a 1% chance of success with the first treatment, and if it fails, you will die. The second treatment will undoubtedly succeed. Your doctor chooses the first treatment, and it is successful
 - Actualism makes an **ex post judgement** based on what the actual outcome is
 - Probabilism makes an **ex ante judgement** by considering the probabilities of each outcome, and evaluating actions based on their **expected value**
 - Able to circumvent moral luck
 - Intuitive because moral decisions take place prior to the consequences; we think of moral decisions to be within the control of the **moral agent**
 - The actualist criterion of rightness does not hold water against our intuitions (see doctor choosing cure thought experiment)thought experiment)
 - Needs to be rooted in the agent's epistemic standpoint (Jackson)
- Sidgwick's conception of Subjective and Objective Rightness
 - From the objective point of view (actualist), the doctor did the right thing since the action was what actually maximised overall happiness
 - From the subjective point of view (tethered to what humans can actually foresee), she had no good reason for thinking that her action would succeed, and thus would be open to criticism from the perspective of subjective rightness
 - The assessment is made in relation to her standpoint, rather than some overarching moral judge
 - Objective rightness consists in the maximisation of overall happiness, while subjective rightness consists in the maximisation of expected overall happiness
 - These theories are concerned with the criterion of the right action and not its decision making procedure
 - But even the criterion of the right action can adopt probabilism (what all this shows is that the decision making procedure CANNOT adopt actualism and subjective rightness because it is beyond the agent's epistemic limits)

- Subjective rightness gives an answer to the infinite horizons objection to utilitarianism
 - Objection: effects of an action stretch indefinitely into the future and we can never know what the consequences of any action we perform will be → infinitely extending causal array of effects
 - So we can never know how to act on the actualist interpretation (or we can never make a judgement on the actualist interpretation)
 - However, probabilism implies that those extending effects can be considered negligible given its low probabilities and unexpected good and bad consequences neutralise one another
 - The direct effects which have significant probabilities are then accounted for → cannot defend oneself for punching someone by arguing for far off benefits

Act Utilitarianism

Standard Act utilitarianism

- Criterion of rightness: An action is right if and only if it produces as much welfare, hedonistically conceived, overall – where everyone's welfare counts and counts equally – as any other action open to the agent
- Decision making procedure: At every point in time t , an agent is to calculate the optimal action that produces the most welfare, by taking into account each and every action and evaluating their long run outcome

Objections to Act Utilitarianism

- Epistemic objection (2 strands of cognitive constraints)
 - Infinite choices: there are infinite variations of each action and infinite different actions to make → implausible for the agent, humanly conceived, to be able to assess the utility maximising action within finite time
 - Infinite time: the consequences of each action are infinitely extending (they can set into motion a chain of events). Not only that, each action would likely cause different events to happen, which in turn cause even more events to happen. The effects of an action are fundamentally multiplicative, as if an infinitely extending causal array. So the actual long run consequences of an action is unknowable to the human
- Self defeating (continuous from the epistemic objection)
 - Even if we accept that it is *possible* to know which action is optimal, one would have to concede that it would be extraordinarily difficult, cognitively demanding and time consuming to reach an accurate and correct conclusion
 - So if we internalise the cost of the decision making procedure into the calculus of welfare itself (since making a decision is itself an action), it seems

hardly plausible that the world with the greatest utility would be one that requires every agent to deliberate in the AU manner

Rule Utilitarianism

Conformity Rule Utilitarianism

- Conformity RU imagines a world where everyone follows the established system of rules
 - Not only is it costless in helping people understand the rules, but it also assumes restlessness in people carrying out the rules
 - Assumes that perhaps people are machines dutifully carrying out the rules
- Problems
 - Collapses into Act utilitarianism
 - We cannot internalise the cost of the decision-making procedure and therefore the optimal decision procedure would be to add as many exception clauses as possible (do not lie unless it prevents a disaster) until it reaches “do not lie unless lying produces greater utility” which essentially means, do the action which produces the greatest utility
 - This is because conformity RU ignores the costs of people having to learn, understand and implement the set of rules in their daily lives
 - There are some rules (such as retaliation rules) where the world would be better off if not complied with
 - The rule that prescribes that you retaliate against attackers: it is optimal for people not to ever have to comply with the rule by not attacking one another in the first place → so there are empty rules that are best not to be complied with
 - With 100% compliance, there would be no place for rules of punishment
 - Rule consequentialism cannot generate or justify rules about how to deter murder, rape... by unmitigated amoralists, unless rule-consequentialism picks its rules with reference to an imagined world where there is internalisation of the envisaged rules by <100% of the population
 - Morality cannot allow for me to bear the cost of the rule by following it and let others benefit from it by not complying with it yet reaping the benefits of general compliance
 - <100% compliance (acceptance) allows for contingencies in rules: Imagine 2 children starving and the rule is you and your friend are obliged to save one each, but your friend does not save him
 - Imagining an imperfect world gives us the space to add in the exception “if others don’t help, do more than what is required of you to prevent the disaster”

- We could reach an optimific rule: People should help others in great need when they could do so at a modest cost to themselves, cost being assessed aggregately, not iteratively (Cullity 1995)
 - Respects the cumulative nature of the cost of helping others as potentially too costly if iteratively assessed

Acceptance Rule Utilitarianism

- Acceptance RU imagines a world where everyone accepts these rules but does not necessarily comply.
 - It does not immediately assume society's seamless ability to comply, and instead begins from the point of society learning these rules.
 - Therefore, the acceptance form of RU makes explicit the rules' costs of internalisation
- It recognises that rules themselves have costs, and not just benefits from the actions they sanction.
 - As rules become more complex, they become more difficult to learn and act upon, incurring large costs onto society
 - So the optimal situation would be for a set of rules that is dissimilar to the act utilitarian decision procedure (though this is an empirical question)

Arguments for Rule Utilitarianism

- Provides a superior decision making procedure
 - Act utilitarianism does not (though this remains to be seen what is the utility maximising world once costs of decision making are internalised – an open empirical question)
 - Act Utilitarians make an explanatory object to Rule Utilitarianism by arguing that it misidentifies what makes actions right or wrong (it is not the violation of rules but the non-maximising of utility)
- Contractualist premise rather than a consequentialist one
 - Contractualist premise is that rules are morally justified just if they are ones to which everyone has good nonmoral reasons to agree
 - The argument then proposes that an act is morally permissible if allowed by morally justified rules
 - Then the argument contents that the only rules that everyone has good nonmoral reason to agree to are the rules whose widespread acceptance would maximise expected value impartially
 - This may not be the case (Rawls argues that people would adopt the maximin principle in such a contractual scenario instead)
 - 'Good nonmoral reason' is used in the sense of maximising welfare
- Taps on intuitively plausible ideas

- What if everyone felt free to do what you are doing → contractually agreed to
- Should we not try to follow a code that could plausibly maximise welfare?
- Avoids the breaking of rules for marginal gains
 - Act utilitarianism would allow for a breaking of promises even if the utility gain is marginal → counterintuitive
 - But we would also believe that breaking promises for a large gain in utility would be justified
 - Act utilitarianism and absolutism (the view that some acts are absolutely right or wrong with no exception) are on opposite ends of the spectrums since AU proposes that actions are only conditionally wrong
 - Rule utilitarianism advantage is that it coheres with when we intuitively feel we can and cannot break rules for the sake of the overall good
 - This is due to the fact that we can only have disaster prevention exception clauses and not marginal gain clauses → general internalisation would be optimal
 - Perhaps rule utilitarianism is self-standing insofar as it only refers to the utility principles as its justification and meta-selection principle and not in its decision procedure and its criterion of rightness
- Avoids the demandingness objection
 - Act Utilitarianism would require us to sacrifice much of ourselves for others → our singular enjoyment <<<< than others suffering
 - But morally unreasonable to demand such level of self-sacrifice → perpetual self-improvement > than what morality demands → perhaps imply supererogatory
 - Imagine allotting offices → you are a senior and you get a corner office but someone benefits more than you and you are indifferent. Even if you do not give up your office for him, no one will say you are morally condemnable → because he has no right to your office in the first place
 - Morality does not require us to sacrifice ourselves normally
 - No space for personal prerogatives
 - But note that RU may still end up being extremely demanding – we just do not know
- Avoids too permissive objection (continuous with the integrity objection)
 - Rule-utilitarian would argue that the best system of rules will incorporate important constraints
 - Would not permit doctors to cut up healthy patients and distribute their organs to save others → because in the aggregate, terrible consequences
 - But this remains an open question

- Motive Utilitarianism: Act Utilitarianism's decision procedure is not simply suboptimal in committing us to overly time consuming welfare calculations, but also **provides us with suboptimal motivation structures which subtract from our overall welfare** (Robert M. Adams)
 - Adams proposes that the motive of the the action itself contributes to the outcome of the action in affecting our psychological states
 - Note that this can extend to any form of utilitarianism: pleasure, desire, OLT
 - Appreciating beauty may be a value in the OLT, we may gain more pleasure from intentionally understanding the paintings, and it may be our desire to revel in art
 - I think Adams' point is this: that **adopting the motive of maximising utility** as per AU **does not provide us with the utility maximising motive structure**, because it takes us away from enjoying the activities we care about or could care about
 - Example: Jack is studying sculpture on the stone choir screen → enjoying this less than other parts of the cathedral and will not remember it well → more happiness on balance if he left earlier
 - This action would not be advised by the AU decision procedure and may be considered wrong if we do not internalise the utility effects of the motive
 - But it may be the case that his caring more about seeing the cathedral than about maximising utility has augmented utility, through enhancing his enjoyment, by more than it has diminished utility through leading him to spend too much time at Chartres → thus he did right by motive utilitarian standards → the motive enhanced his enjoyment because he was not looking at the cathedral for some instrumental end but rather for its intrinsic utility and thus could appreciate it better
 - He may desire to enjoy the paintings; but if he is only focused on maximally fulfilling his desires, would **this not conflict with his ability to actually enjoy the art around him?** And for pleasure, would the **perpetual concern** for utility maximisation **not subtract away from our overall pleasure?**
 - Utilitarianism then cannot simply evaluate utility on the basis of **what action** is committed, but must internalise the utility effects of **how** the action is committed
 - So Adams extends the point that AU's decision making is not utility maximising: it is not utility maximising in also how we go about conducting our lives → being rational calculators of welfare itself undermines our ability to experience pleasure (in the things around us) or truly satisfy our desires

- This is beyond the objection that it consumes too much time and thus cost us much
 - So in internalising the utility effects of the motivation structure, Adams concludes that the AU motivation structure is not utility maximising
 - We can see how RU can plausibly **incorporate having the appropriate motivation structure as a rule**
 - This shows that RU provides greater flexibility in our decision procedures than AU
 - Disagree with Adams on the incompatibility between motive utilitarianism and AU: motive utilitarianism is not incompatible with AU's explanatory component/criterion of rightness of what actions are right – we just need to extend the analysis of actions to the dimension of how they are committed/their motive since it has utility effects
 - The second key insight from Adam's is that RU's decision procedure is not necessarily utility maximising because it retains an inappropriate motivation structure; so **MLAU does not necessitate rules as the decision procedure**
 - RU prompts people to be **motivated to act in a certain way because of the established rule** and not for the sake of the action itself
 - Motive utilitarianism requires Jack to adopt a motivation structure that inherently appreciates the value of revelling in the art around him; it is a motivational structure that compels an agent to act for action X because **he appreciates the value of action X**
 - The extreme reading is that any motivational structure other than acting for some action X itself incurs some costs in the decision procedure (of course the intensity varies); so we need to internalise the cost of the suboptimal motivation structure: makes the question of whether MLAU incorporating RU as decision procedure is $\max(u)$ even more of an open question
- Congruent with the universalisation test
 - If your friend tells you that it is fine for him to evade taxes or cheat in exams because the deviation from generally upheld norms will not cause any bad consequences, your reaction will likely be 'What if everyone thought it would be OK to do that?'
 - The moral rules are universally applied and upheld → static in nature
 - The moral imperative does not change based on what other people do → when they don't help someone, it does not force you to help more so it avoids the too demanding charge
 - But we also find it counterintuitive if our moral rules are constantly in flux

Objections to Rule Utilitarianism

- Criterion of rightness incoherent with the principle of utility
 - Principle of utility: about the individually utility maximising action → so RU deviates from this completely by pegging rightness to conforming to the set of rules
 - Solution: suggest that the rule-based criterion of rightness still explains rightness and wrongness on the basis of maximising utility, and that it is actually more coherent with achieving a history/world with the greatest utility since it can account for the **interaction of actions across time and between people**
 - Crisp's edit of the principle of utility: The best actions are those which feature in the best possible history
 - Decompose the notion of history: not just about actions across time, but between people as well
 - So, the principle of utility sanctions a set of actions that collectively maximises overall utility, rather than individually utility-maximising actions → AU's decision criteria commits the fallacy of composition (see W3 essay and extend the argument to across time)
 - Revised Objection: even if we accept Crisp's suggestion, RU assumes that the in set of actions that collectively maximises utility would include rules of actions that apply generally
 - RU assumes that, usually, the set of actions that collectively $\max(u)$ implies that we must always not lie since we create a general rule dictating whether or not it is acceptable to do an action
 - But intuitively, the maximising set of actions would have deviations from the rule across time and between people (in some cases, it would be best not to lie and in others it would be best to lie)
 - Even incorporating contingencies may be insufficient; in speaking about the criterion of rightness, we can ignore the costs of decision making
 - So it is hard to accept that the RU criterion of rightness adequately captures the principle of utility
- No unique best system of rules
 - Implicitly assumes that there would be one unique best system of rules
 - And the Utilitarian principle supposes indifferences between two actions (and by extension, sets of rules) that are maximally and thus equally conducive of welfare
 - But this indifference cannot hold since everyone needs to adopt a single system of rules

- Consider 2 systems that one differs only in asking you to drive left or right → not implausible to fill in the details so that these systems are equally good
 - Utilitarianism here is indifferent between you following a set of rules that differ in asking you to drive left or right
 - So people drive on both sides simultaneously and crash and de
- But this objection is mistaken → rule utilitarianism would **only be indifferent between the set of rules at the point of choosing the set of rules**
 - It would not allow for one to implement both set of rules simultaneously because it is obvious that such simultaneous implementation would be mistaken
- Assumed existence of a best set of rules distinct from AU
 - Remains an open empirical question → not a denial of RU, but cautioning how much progress RU actually makes over AU
 - It may be the case that even after internalising the cost of decision procedures, AU's decision procedure is still the best
- Problem with partial compliance/acceptance
 - Require rules that handle people who transgress rules even though majority accept the rules (and some do not comply)
 - Idealistic: If everyone accepted my pacifist principles, there would be no wars and no need to take up arms. Therefore, a rule about not taking up arms should be established
 - Non-problem: obviously just incorporate rules on punishment conducive to the overall welfare → compel people to comply too
- Disaster prevention
- Incoherence with common sense in why actions are wrong
 - Fundamental disagreement about what makes actions right or wrong (justification)
 - Rule Utilitarianism: what makes actions wrong is that fact that this action is forbidden by a system of rules whose general acceptance would bring about the greatest total of well-being
 - Common sense: what makes it wrong is something about you and the person you made a promise to is simply the fact that you made a promise to him
 - Awkward: people do not follow rules for the sake of following rules, but for some underlying reason → so it attacks RU's criterion of rightness for assessing rightness from the standpoint of following rules itself
 - Solution: suggest that the rule-based criterion of rightness still explains rightness and wrongness on the basis of maximising utility, and that it is actually more coherent with achieving a history/world with the greatest utility since it can account for the **interaction of actions across time and between people**

The Collapse Argument (2 strands)

- First strand: Rule Utilitarianism collapses into Act Utilitarianism
 - If what is morally right is the action that produces the maximally good outcome, then the set of rules would eventually have to specify infinite contingencies until the decision procedure collapses into “Do X unless it does not maximise the good”
 - But obviously we can tackle this using acceptance Rule Utilitarianism
- Second strand: Rule Utilitarianism collapses into customary morality
 - If you do not include exceptions, then you are not a consequentialist because you do not truly use utility/outcome values in your analysis
 - Following rules are not necessarily consequentialist; the justification of these rules must be consequentialist (that they produce the greatest utility) → so RU threatens to depart from utilitarianism itself in its explanatory aspect
 - Solution: we will probably have to include some exceptions
 - Our utility maximisation problem would be to maximise utility with respect to the complexity of the rules such that the net utility from “utility from a world following the rules - the cost of internalisation” is maximised
 - This implies that it is not necessarily the case that we have simple rules of thumb as in customary morality
 - RU is still grounded on maximising utility – it just needs to take into account the costs of having complex rules
 - Moreover, by needing to consult the Act Utilitarian principle (which is itself a rule) during times of conflicting rules, it shows that the fundamental principle is grounded on utility and maximising outcome values
- Brad Hooker’s conjecture: we can cleanly avoid the collapse argument by not beginning from an overriding commitment to maximise utility
- Bykvist Objection of the Collapse argument
 - Bykvist asks us to imagine a simultaneous single-shot game

A/B	B go to Oz	B go to Kings
A go to Oz	6	0
A go to Kings	0	10

- If we both go to Oz, that would be right according to act utilitarianism
 - This is because act utilitarianism takes the other’s action (going to Oz) as a given – exogenously determined

- So AU evaluates an agent's action on whether it maximises utility in a given scenario, even when the factors are beyond the agent's epistemic standpoint
- And we want to say they are both wrong! (but Bacceli contests this)
- The point is that acting according to act utilitarianism will not produce the best circumstances in itself → but this is not acting according to act utilitarianism → the evaluation is ex post
 - RU on the other hand would imagine the set of actions that **collectively determine** the optimal outcome (does not take any action as exogenously determined since RU follows a universalisation test)
- Perhaps we need to redefine the best history as one that is in relation only to the current agent's actions (so subject to exogenously determined actions)
 - After all, the right action intuitively is influenced by other player's actions → like donating, or saving the child

Multi-level Act Utilitarianism

(Mill's) Multi-level Act Utilitarianism

First level of moral thinking	<p>Act utilitarian principle as the sole criterion of moral rightness, and the justification of all actions</p> <p>Actions are wrong because they contravene this principle, and not the rules formed from it</p>
Second level of moral thinking	<p>Application of rules in the form of a decision procedure for what to do (customary morality), which itself is founded on the act utilitarian principle to maximise utility → prevent calculations which will defeat the goal of utilitarianism itself</p>
Intermediate between second and third level	<p>Conflicts between secondary rules are arbitrated by the act utilitarian principle</p>
Third level of moral thinking (useless)	<p>Philosophy itself: this is engaging in discourse in the sense of understanding morality and 'upgrading' it via metaethics</p>

Reflections on Multi-level Act Utilitarianism

- An actualist conception of the criterion of rightness seems to imply that no decision

procedure can truly achieve maximal utility/or at least an extraordinarily difficult empirical question

- A decision procedure requires us to either (1) calculate at every point in time what is the best action from the set of all possible actions or (2) specify conditions under which a specific class of actions (not lying) is the best one
 - So (2) implies that our rule-based decision procedure can be quite sophisticated: Do not lie (because it is utility maximising) unless it prevents disaster → contingency clauses
 - (2) also implies that there exists a pattern regarding the utility-conduciveness of our action: under conditions X and Y, action Z would be utility maximising
 - This is not obvious at all → randomness and contingent factors
- So, given that the utility-conduciveness of our actions are unlikely to follow a trend such that it is amenable to prediction, if we look at utility-maximisation solely as a question of achieving the 'actions featured in the best possible history', then MLAU is unlikely to lead us to the right action on many scenarios (is this then a good decision procedure?)
- Clearly, in conceiving of a decision procedure, we ought to internalise the costs of the decision procedure itself
 - Note that the decision procedure incurs cost for every decision made (so cumulatively it is significant)
- So now we can set up the utility maximisation problem
 - $U(x) = W(x) - C(x)$
 - W is solely the utility of the 'history' achieved independent of decision making costs; it is increasing in x
 - X is the level of complexity in how we make our decisions
 - C is the cost of the decision making procedure; it is increasing in x
 - The maximisation problem is then $\max(U(x) = W(x) - C(x))$ and not $\max(U(x) = W(x))$
 - It follows that the best possible history $W(x^*)$ where x^* represents infinite complexity would likely not imply maximal utility
 - The range of $W(x)$ likely has a supremum – a maximum utility given we all optimise our actions (abstracting away decision costs)
 - It is strange to think of the best possible history having infinite utility – instead, we think of it as having the most (finite) utility
 - On the contrary, $C(x)$ is unlikely to have a supremum
 - Because we need to account for an (i) infinite number

of outcomes over (ii) an infinite number of decision points, we can easily see that as x tends to infinity, it is obvious that $C(x)$ will tend to infinity too

- So it cannot be the case that the maximal utility situation is to have maximal complexity
 - So MLAU or Rule Utilitarianism as a decision procedure will not collapse into AU
 - On the contrary, AU would collapse into MLAU
- Probabilistic criterion of rightness (simpler argument)
 - However, if we take the criterion of rightness to be referring to an action that maximises the expected value (see Jackson's argument) then MLAU would likely lead to the best outcome
 - Certain actions (like not lying unless under exceptional circumstances) have the high probabilities of having a good outcome, and conversely low probabilities that it turns out to be the negative outcome
 - So it makes sense to develop rules as a decision making procedure to capture their generally higher probability of achieving a good outcome
 - There are some actions (like lying) that are generally riskier in
 - In fact, the expected value of an action (Probability of outcome x \times value of outcome for all outcomes given action i), the expected value criterion of rightness is especially suited for a rule-based decision procedure
 - The probability of an outcome given an action i is likely assessed from a general observation of the distribution of outcomes for any given action
 - In this sense, a high probability of an outcome A for action i implies that in general, action i is likely to lead to A
 - It follows that, ceteris paribus, the probability of an outcome A for action i in most normal circumstances is unlikely to (or should not) change \rightarrow so it is more amenable to using a rule-based decision procedure to achieve the right action
 - This is unlike actualism where probabilistic randomness means that what is deemed as the right action would fluctuate constantly and would be unpredictable
- Bite the bullet
 - The main tension is how can the decision procedure (RU) and the criterion of rightness deviate so much? How can the decision procedure not lead to the right action 100% of the time
 - Answer: it leads to the right action most of the time and so it is reliable \rightarrow ask ourselves what makes for a good decision procedure?
 - A decision procedure does not need to be always correct, merely reliable such that people can use it

- It is like a GPS – sometimes it leads you to the wrong place, but 80% of the time it leads you to the correct place
- Moreover, the RU decision procedure is likely only to fail in minor cases: i.e. white lying that you scored badly in an exam to make your friend feel better and not disaster cases like lying to the murderer looking for your friend